

Explanation, Pseudocode, Data Structures

High Level Explanation:

The algorithm is centered around how the controller will determine where the given switch should forward the packet to. The controller uses a data structure to maintain a state and to incrementally learn about the topology. My approach was to ensure the controller could have a table for each switch. When a switch sent a packet to the controller, it would first update/add the source information from the packet (via a source to port mapping). Then it would attempt to deal with the destination information the packet specified (finding an entry that has the destination address its corresponding port can be sent back to the switch). However, it was essential this search for the dest-port mapping entry was only done in the switches own table and not in all the switch tables the controller has, which is why each element in the controllers data structure is representing a different switches table. This left only two possible cases to handle: either the switch table had the matching dest-port mapping entry, or it didn't have the entry. If a match was found then we simply send the corresponding port value back to the switch and if no match was found then we tell the switch to flood all of its other ports.

Thought Process Through the Algorithm:

Steps one and two are about making sure the controller has a table for the switch that sent the packet and adding the source-related mapping to the table. This needs to be done before any destination-related logic is dealt with.

STEP 1

Check if the controller already has a switch-table for the switch which sent the packet: if not, then create empty table for this switch

STEP 2

Use the packets source address and 'packetInPort' to add the address-port mapping of the source as an entry in the switches table (unless the table already has an entry for this source address. (the source address will be the 'value' of the 'mac' key, so we would be iterating through the elements in the switch_table, and comparing the 'value' of the dictionaries first elem, to packet.src (since the first elem in the dictionary is the 'mac': 'address' key-value pair, and we want to check 'address', which is the associated value)

STEP 3

Now the controller will take the packets destination address and look for a matching entry in the switch_table (using the same iteration that was done to check for a source address entry in step 2).

CASE 1: MATCH FOUND --> get 'port' from match and put it in the message back to the switch

CASE 2: NOT FOUND: --> flood to all other of the switches ports.

More detailed explanations of the two cases:

CASE 1: If a matching entry was found, then we can get the corresponding 'port' which is the important piece of information the switch was asking for when it sent the packet to the controller in the first place. We can then send this port value in the message being sent back to the switch and it will then know where to forward the packet to.

CASE 2: If a match was not found in the switches table, then we send a message back to the switch which tells it to flood all of its 'other' ports. Because of the initial steps of the algorithm, when this flood eventually returns the needed host-port information, it will be entered into the controllers topology right away instead of needing to wait for the destination-related address to be dealt with. This is what allows for the topology to be iteratively mapped out, even when starting out with no knowledge of the host/switch relationships.

PSEUDO-CODE (Cormen):

```
if switchID not in all_switch_tables:  
    add empty table for this switchID to all_switch_tables  
  
switch_table = all_switch_tables[switchID] // table of the switch that sent the packet  
if packet.src not in switch_table: // add the source-port mapping to the switch table  
    switch_table.append(packet.src, inputPort)  
  
// search the table for entry matching packet.dst...  
if packet.dst in switch_table: // entry found  
    foundPort = switch_table[packet.dst]  
    msg = (foundPort)  
    sendBackToSwitch(msg)  
else: // entry not found...  
    msg = (flood)  
    sendBackToSwitch(msg)
```

Data structure explanation: (also in ethernet-learning.py)

```
36 # ****  
37 # DATA STRUCTURE OVERVIEW:  
38 # ****  
39  
40 # CONTROLLER:  
41 # the controller learns/maps the network with a DICT called all_switch_tables. The key-value pairs of this dict are the switchID and the corresponding switch table  
42 # (key == switchID, val == switch_table)  
43 # This allows the controller to quickly check if it already has a switch_table for the switch that just sent it the packet in question, while also ensuring it doesn't have  
44 # multiple tables for the same switch.  
45  
46 # SWITCH-TABLE:  
47 # a switch table is a LIST, called switch_table, where each elem is itself a DICT (and this DICT will have TWO elements, each of which is a key-val pair)  
48 # The 2-elem DICT's will be as follows: (one element is for the mac address and the other is for the port)  
49 # left elem: key=mac:addr-val  
50 # right elem: key=port:port-val  
51  
52 # The primary task the controller needs to be able to do is search through a particular switches table to match the packets address (either source or dest) with the address in the table.  
53 # Because each element of the table is a dictionary, the address the controller needs in the table is the 'value' of the first element.  
54 # So for each entry in the table, it will evaluate that entries first element by referencing the "mac" key to get the associated 'value', which is the actual address for the entry, and this  
55 # is the value the controller will be trying to match the packets addresses with.  
56
```

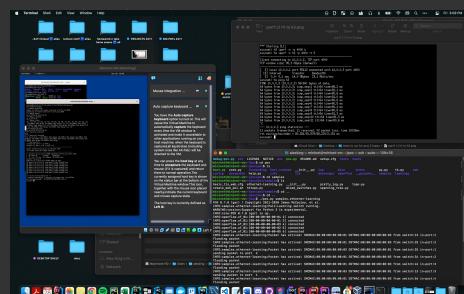
Part 3 Results for each topology:

topology a

L1

BW (Mbit/sec) : 19.3

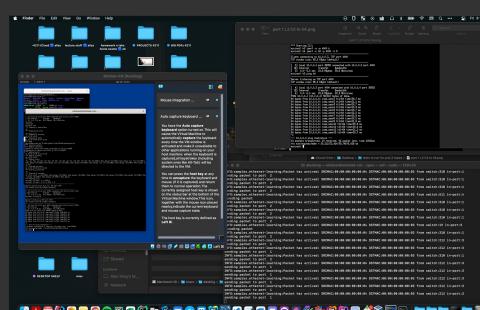
Latency(ms) : 80.6



L2

BW (Mbit/sec) : 37.0

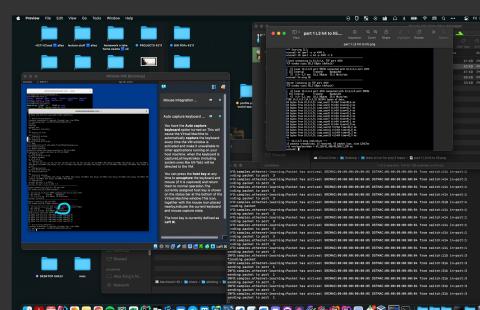
Latency(ms) : 21.6



L3

BW (Mbit/sec) : 25.0

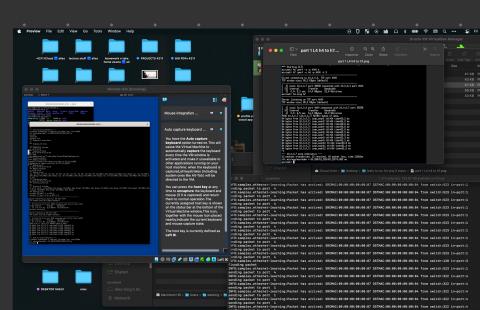
Latency(ms) : 41.3



L4

BW (Mbit/sec) : 17.6

Latency(ms) : 60.9



topology b

LATENCY

level 1 link(s) : 61.0

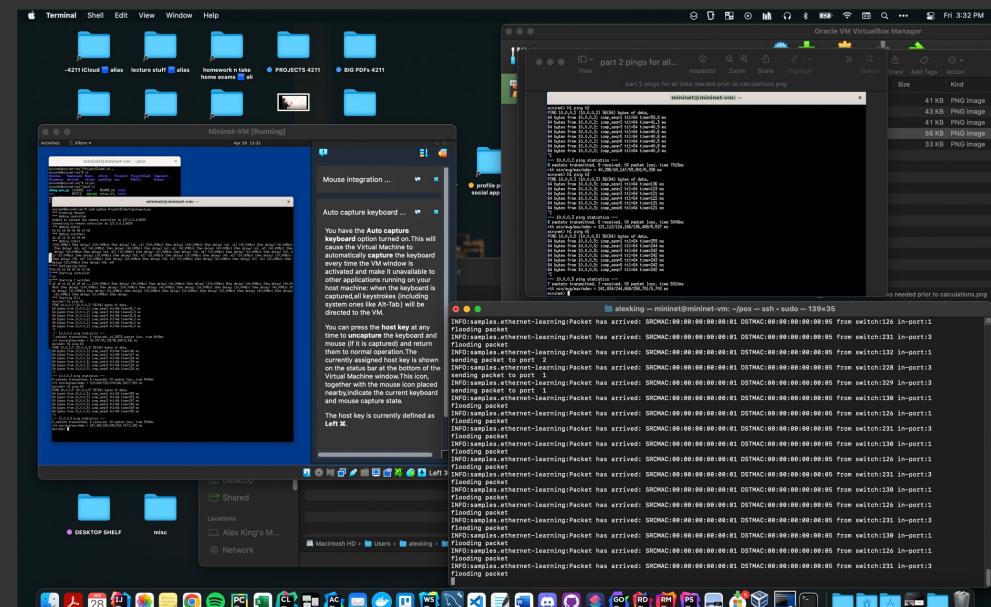
calculator...

total = 250

level 2 link(s) : 41.6

calculator

total = 129.1



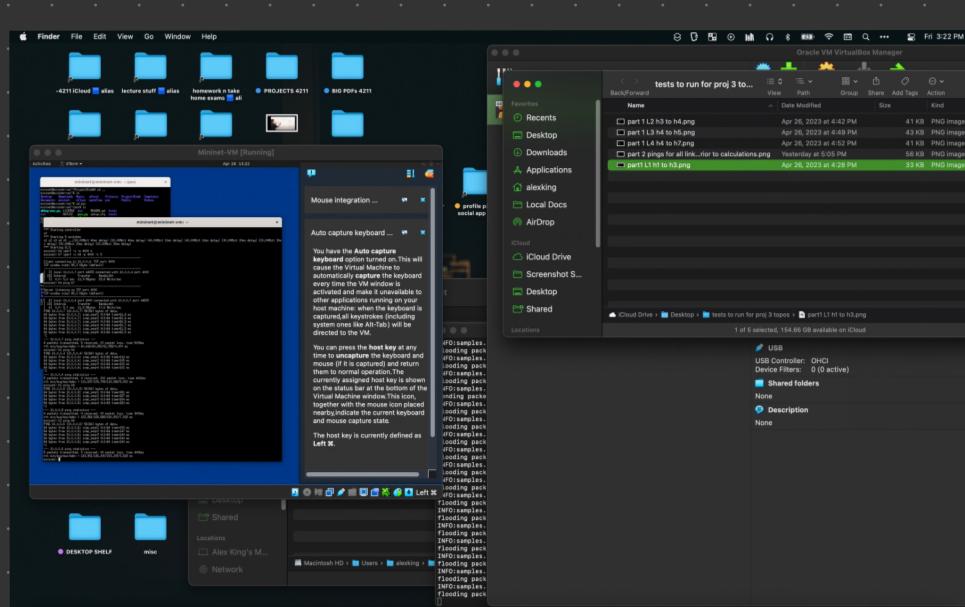
$$\text{Latency}_2 = \frac{\text{total} - 2(\text{Latency}_1)}{2} = \text{Latency}_1$$

$$= \frac{(129.1 - 2(25.0))}{2} = 41.6$$

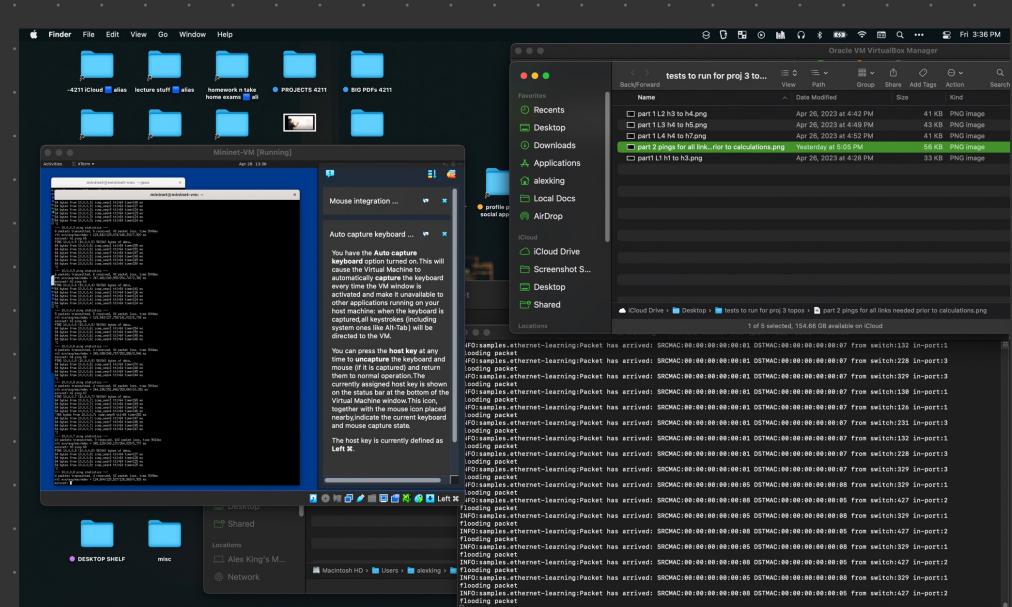
$$\text{Latency}_1 = \frac{\text{total} - (\text{Latency}_2 \times 2)}{2}$$

$$= \frac{(250 - (41.6 \times 2))}{2} = \frac{(250 - 83.2)}{2} = 83.4$$

Extra pings to demonstrate each host can reach any of the other hosts



Extra pings to demonstrate each host can reach any of the other hosts



Comparisons

Comparing part 1 results to part 3 results (topology-a)

PART 1

L1 Bandwidth: 23.1 Mbits/sec
L1 Latency: 81.6 ms

L2 Bandwidth: 39.5 Mbits/sec
L2 Latency: 22.4 ms

L3 Bandwidth: 30.2 Mbits/sec
L3 Latency: 41.5 ms

L4 Bandwidth: 22.6 Mbits/sec
L4 Latency: 61.5 ms

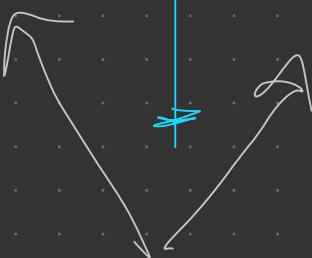
PART 3

B/N (Mbit/sec) : 19.3
Latency(ms) : 80.6

B/N (Mbit/sec) : 37.0
Latency(ms) : 21.6

B/N (Mbit/sec) : 25.0
Latency(ms) : 41.3

B/N (Mbit/sec) : 17.6
Latency(ms) : 60.9



The similar link values show our part-3 algorithm worked for topology-a

Comparing part 2 results to part 3 results (topology-b)

PART 2

LATENCY

Total for level 1 test: 244.9
Level 1 link(s): 60.6

Total for level 2 test: 124.2
Level 2 link(s) : 40.4

Total for level 3 test: 47.6
Level 3 link(s) : 21.6

PART 3

LATENCY

Total for level 1 test: 250
Level 1 link(s): 61.0

Total for level 2 test: 129.1
Level 2 link(s) : 41.6

Total for level 3 test: 45.7
Level 3 link(s) : 22.8



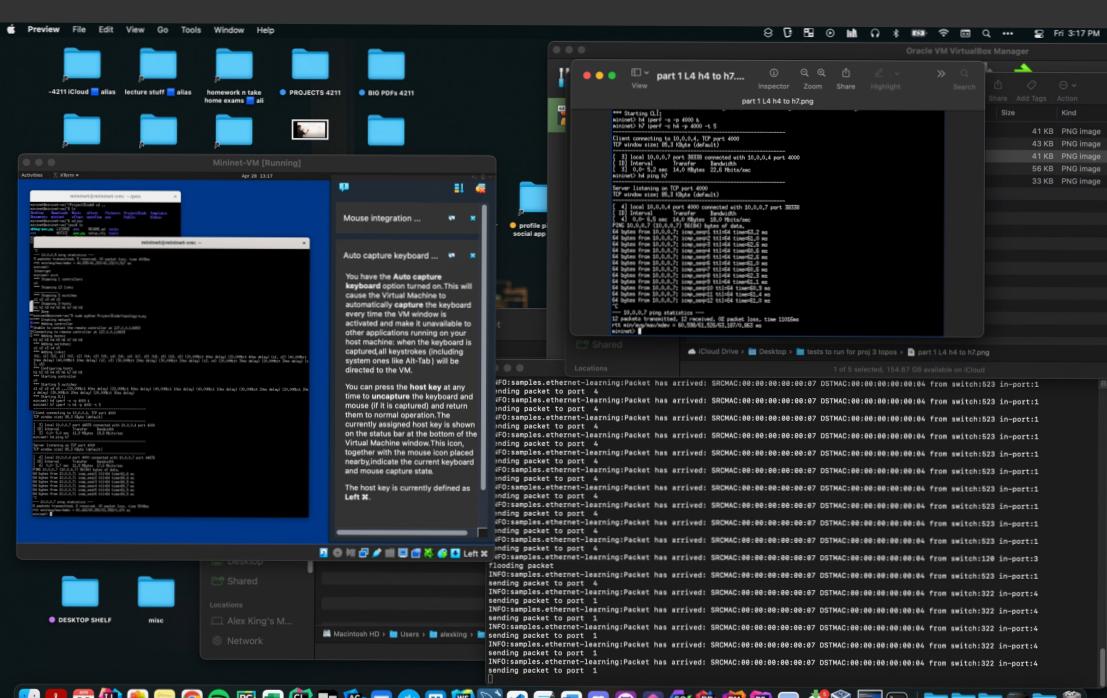
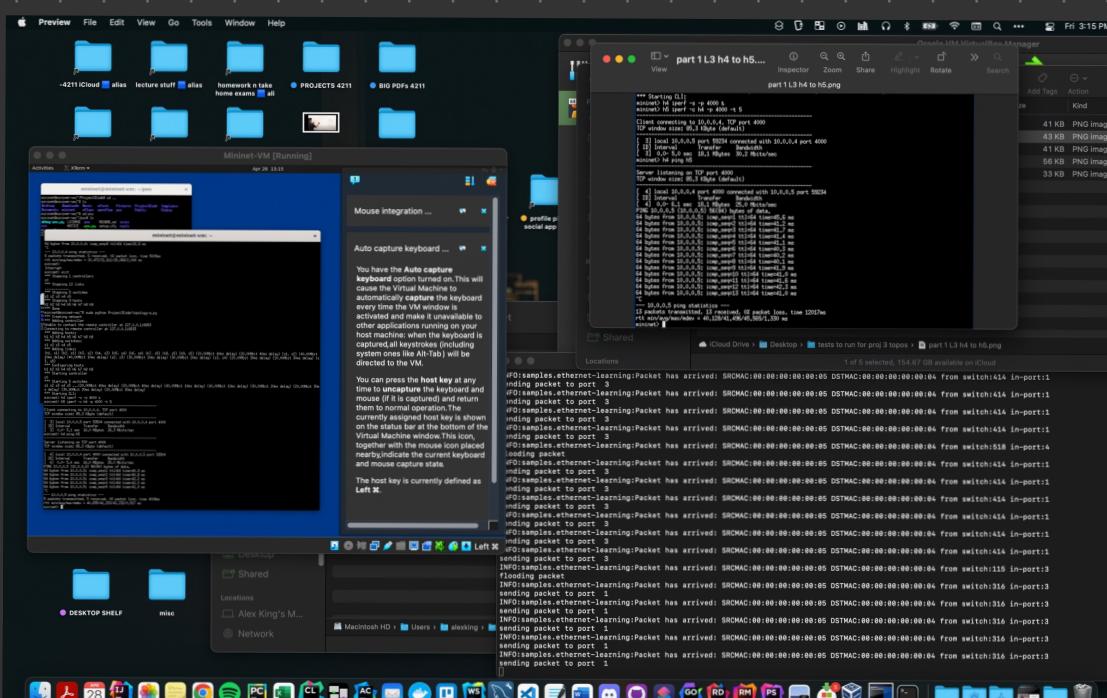
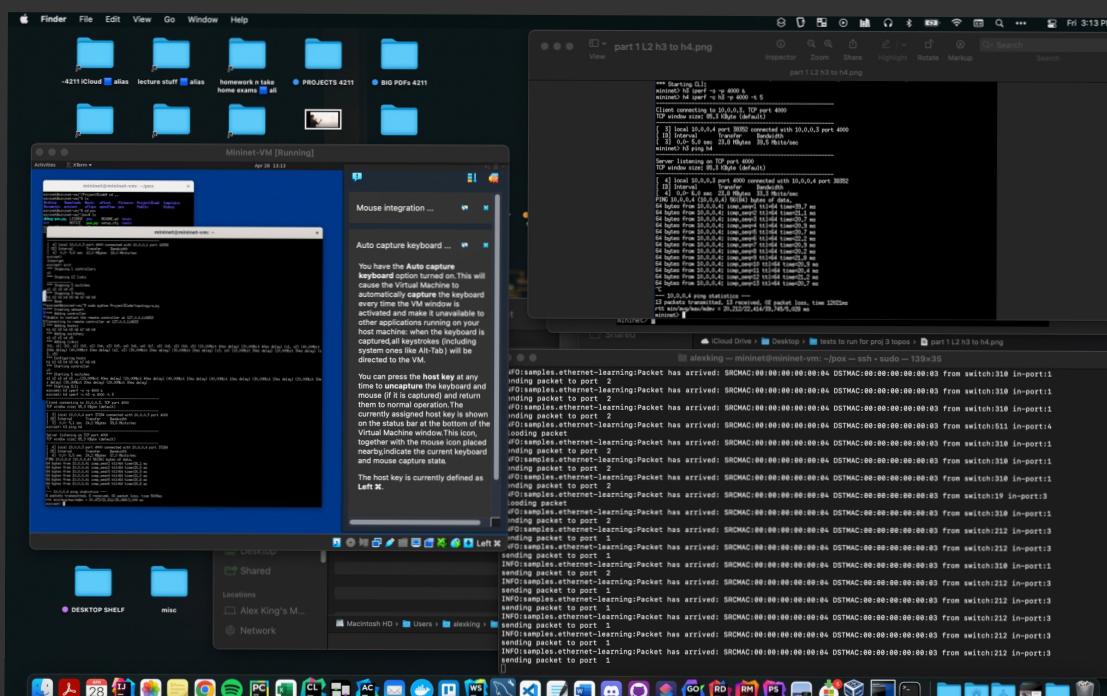
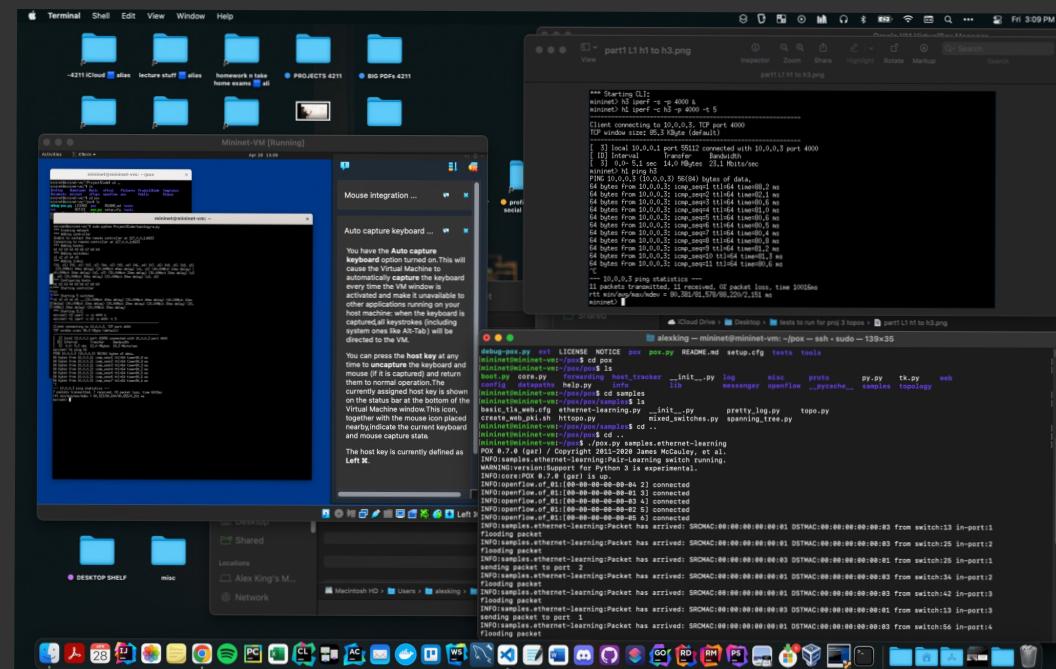
The similar link values show our part-3 algorithm worked for topology-b

```
mininet@mininet-wm: ~
```

```
tcpdump: listening on eth0, link-type ENETTYPE (Ethernet), promiscuous mode off
[...]
Layer 3 links
tcpdump: listening on eth0, link-type ENETTYPE (Ethernet), promiscuous mode off
[...]
Layer 2 links
tcpdump: listening on eth0, link-type ENETTYPE (Ethernet), promiscuous mode off
[...]
Layer 1 links
tcpdump: listening on eth0, link-type ENETTYPE (Ethernet), promiscuous mode off
[...]
```

ENLARGED SCREENSHOTS

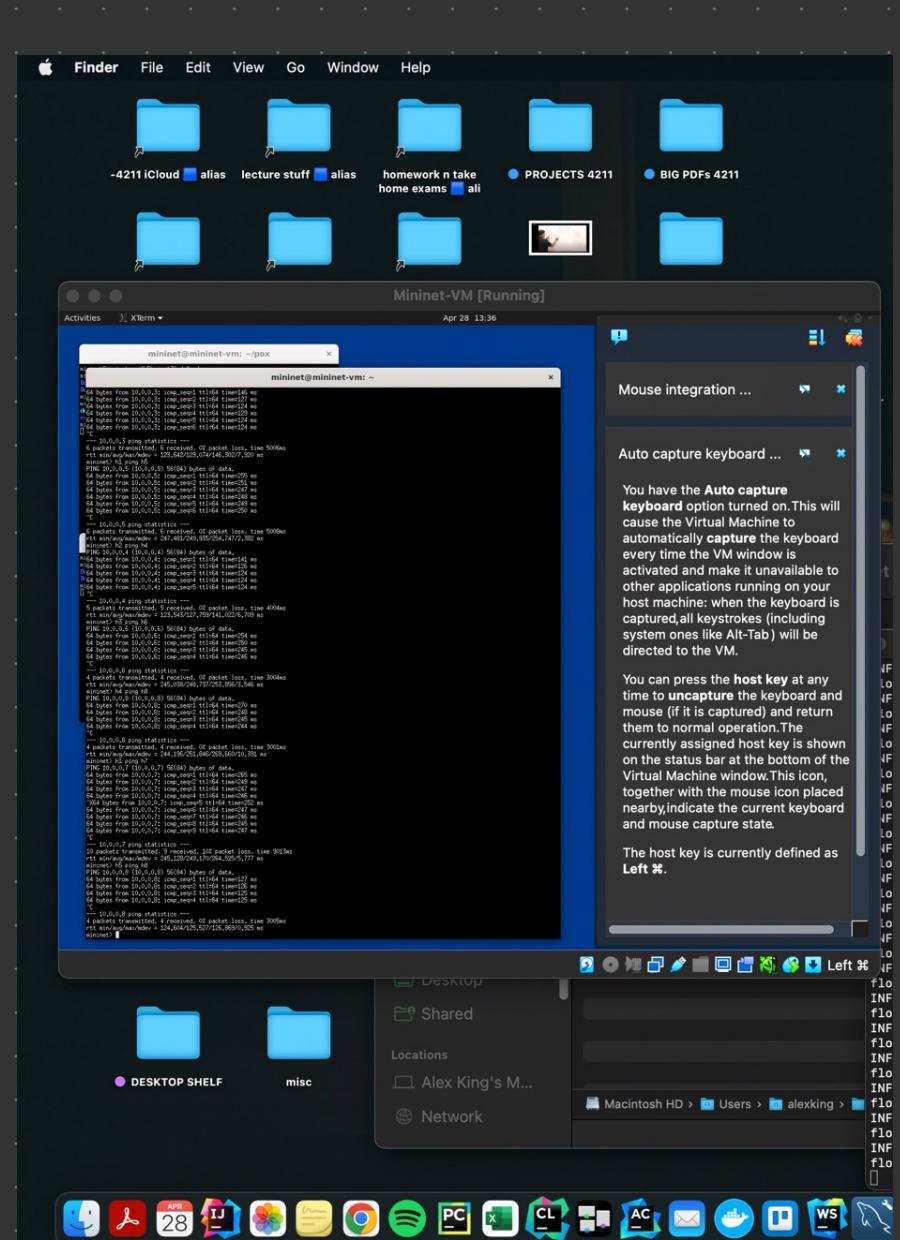
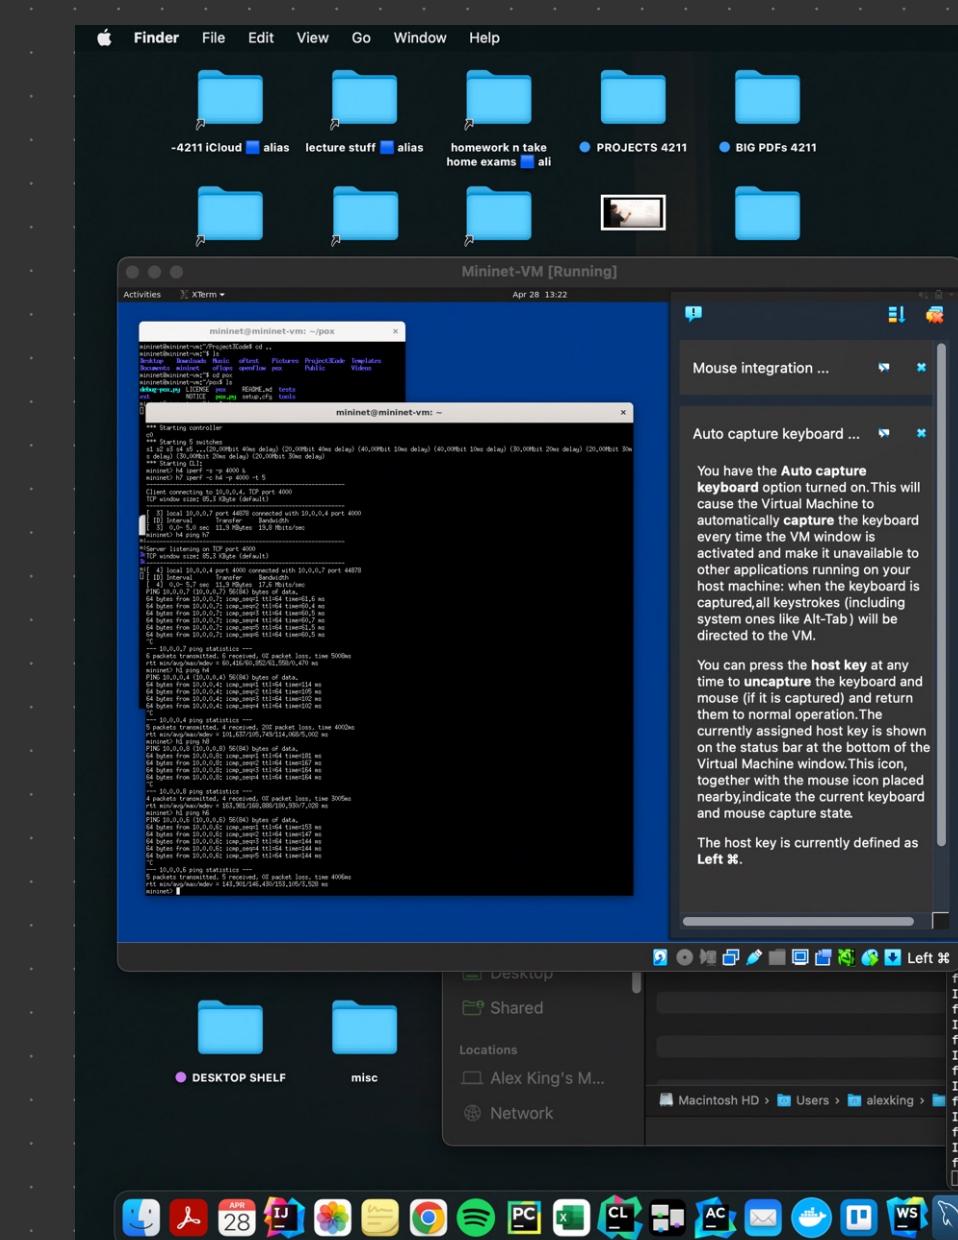
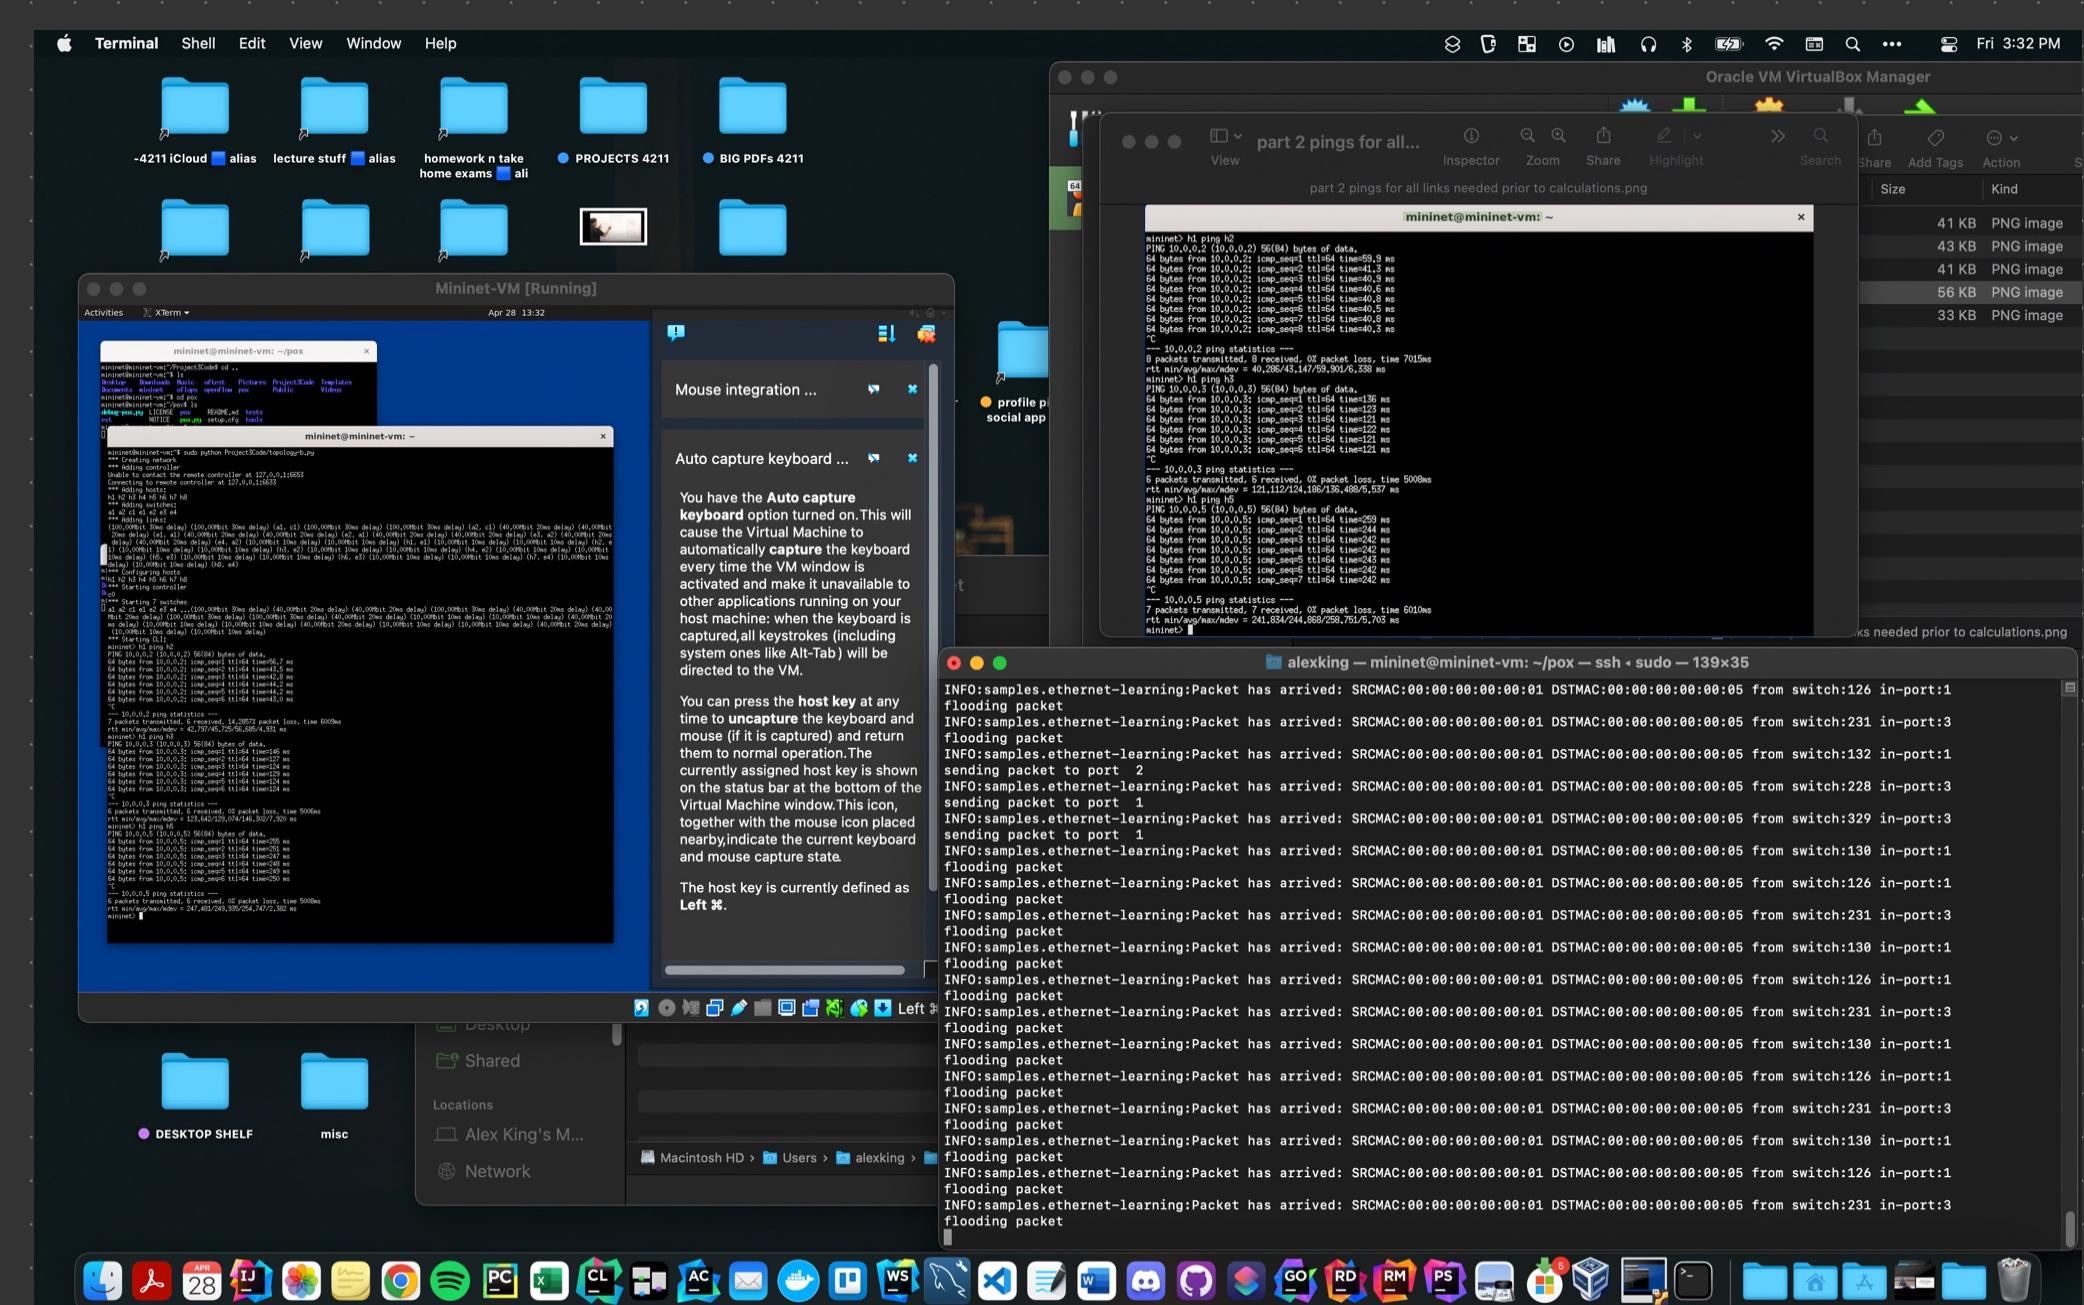
topology a



ENLARGED SCREENSHOTS

topology b

+ Extra PINGS



SCRATCHWORK

Accepts [] [Packet]

OF PackIn

(b/c switch sending it has no entry for dest)

Switch Table (Dictionary)
MAC part

(1) First add mapping (SRC MAC) \rightarrow (Switch Port)

Switch Table
MAC part

h1 Mac	1

(2) Look in table for (DEST MAC) that matches the packet

? ?

Entry found
NOT FOUND

Tell switch to Flood all its ports except

Controller Table

MAC	Switch	Port
h1 Mac	X	
h2 Mac		X

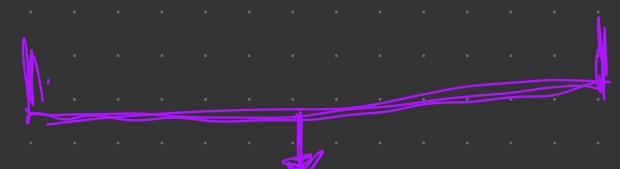
switchTable.append({ Key : value , Key: Value }) Controller Table

Controller Table

MAC	Port
h1 Mac	
h2 Mac	

I elem in the switchTable

switchTable.append({
 ^(mac)Key : value ,
 ^(port)Key: Value })



switchTable = {

:

"

Dict with 2 elements

→ each elem has a keyVal pair